



```
In [1]: import pandas as pd
```

```
In [2]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
```

```
In [15]: df=pd.read_csv("C:\\\\Users\\\\Admin\\\\Desktop\\\\emp.csv")
```

```
In [16]: df.head()
```

```
Out[16]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome
0	41	Yes	Travel_Rarely	1102	Sales	1
1	49	No	Travel_Frequently	279	Research & Development	8
2	37	Yes	Travel_Rarely	1373	Research & Development	2
3	33	No	Travel_Frequently	1392	Research & Development	3
4	27	No	Travel_Rarely	591	Research & Development	2

5 rows × 35 columns

```
In [17]: df.isna().sum()
```

```
Out[17]: Age          0  
Attrition      0  
BusinessTravel  0  
DailyRate       0  
Department      0  
DistanceFromHome 0  
Education        0  
EducationField    0  
EmployeeCount     0  
EmployeeNumber    0  
EnvironmentSatisfaction 0  
Gender          0  
HourlyRate       0  
JobInvolvement    0  
JobLevel          0  
JobRole           0  
JobSatisfaction    0  
MaritalStatus      0  
MonthlyIncome      0  
MonthlyRate        0  
NumCompaniesWorked 0  
Over18            0  
OverTime          0  
PercentSalaryHike 0  
PerformanceRating 0  
RelationshipSatisfaction 0  
StandardHours      0  
StockOptionLevel    0  
TotalWorkingYears   0  
TrainingTimesLastYear 0  
WorkLifeBalance    0  
YearsAtCompany      0  
YearsInCurrentRole 0  
YearsSinceLastPromotion 0  
YearsWithCurrManager 0  
dtype: int64
```

```
In [7]: import pandas as pd  
  
print(df.head())  
print(df.info())  
print(df.isnull().sum())
```

```
Age          0
Attrition    0
BusinessTravel 0
DailyRate     0
Department    0
dtype: int64
<class 'pandas.core.series.Series'>
Index: 35 entries, Age to YearsWithCurrManager
Series name: None
Non-Null Count Dtype
-----
35 non-null   int64
dtypes: int64(1)
memory usage: 1.6+ KB
None
0
```

```
In [18]: df['Attrition'] = df['Attrition'].map({'Yes': 1, 'No': 0})
```

```
In [19]: df
```

```
Out[19]:      Age Attrition BusinessTravel DailyRate Department DistanceFromHome
0       41        1   Travel_Rarely      1102      Sales
1       49        0   Travel_Frequently     279  Research & Development
2       37        1   Travel_Rarely      1373  Research & Development
3       33        0   Travel_Frequently     1392  Research & Development
4       27        0   Travel_Rarely       591  Research & Development
...
1465    36        0   Travel_Frequently     884  Research & Development
1466    39        0   Travel_Rarely       613  Research & Development
1467    27        0   Travel_Rarely       155  Research & Development
1468    49        0   Travel_Frequently     1023      Sales
1469    34        0   Travel_Rarely       628  Research & Development
```

1470 rows × 35 columns

```
In [20]: df['Attrition']
```

```
Out[20]: 0      1
         1      0
         2      1
         3      0
         4      0
        ..
       1465     0
       1466     0
       1467     0
       1468     0
       1469     0
Name: Attrition, Length: 1470, dtype: int64
```

```
In [21]: df = df.drop(['EmployeeCount', 'EmployeeNumber', 'Over18', 'StandardHours'], axis=1)
```

```
In [22]: df
```

```
Out[22]:    Age Attrition BusinessTravel DailyRate Department DistanceFromHome
0      41        1   Travel_Rarely     1102      Sales
1      49        0   Travel_Frequently    279  Research & Development
2      37        1   Travel_Rarely     1373  Research & Development
3      33        0   Travel_Frequently    1392  Research & Development
4      27        0   Travel_Rarely      591  Research & Development
...
1465    36        0   Travel_Frequently    884  Research & Development
1466    39        0   Travel_Rarely      613  Research & Development
1467    27        0   Travel_Rarely      155  Research & Development
1468    49        0   Travel_Frequently    1023      Sales
1469    34        0   Travel_Rarely      628  Research & Development
```

1470 rows × 31 columns

```
In [23]: le = LabelEncoder()

for col in df.columns:
    if df[col].dtype == 'object':
        df[col] = le.fit_transform(df[col])
```

```
In [24]: df
```

```
Out[24]:   Age Attrition BusinessTravel DailyRate Department DistanceFromHome
0    41        1            2      1102          2
1    49        0            1      279           1
2    37        1            2     1373           1
3    33        0            1     1392           1
4    27        0            2      591           1
...
1465  36        0            1      884           1          2
1466  39        0            2      613           1
1467  27        0            2      155           1
1468  49        0            1     1023          2
1469  34        0            2      628           1
```

1470 rows × 31 columns

```
In [25]: X = df.drop('Attrition', axis=1)
y = df['Attrition']
```

```
In [26]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [27]: sc = StandardScaler()
X_train_scaled = sc.fit_transform(X_train)
X_test_scaled = sc.transform(X_test)
```

```
In [28]: log_model = LogisticRegression()
log_model.fit(X_train_scaled, y_train)
log_pred = log_model.predict(X_test_scaled)
```

```
In [29]: tree_model = DecisionTreeClassifier()
tree_model.fit(X_train, y_train)
tree_pred = tree_model.predict(X_test)
```

```
In [33]: from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [34]: print("Logistic Regression Accuracy:", accuracy_score(y_test, log_pred))
print(confusion_matrix(y_test, log_pred))
print(classification_report(y_test, log_pred))
```

```
Logistic Regression Accuracy: 0.8945578231292517
[[249  6]
 [ 25 14]]
      precision    recall   f1-score   support
          0       0.91      0.98      0.94      255
          1       0.70      0.36      0.47       39

   accuracy                           0.89      294
macro avg       0.80      0.67      0.71      294
weighted avg    0.88      0.89      0.88      294
```

```
In [35]: print("Decision Tree Accuracy:", accuracy_score(y_test, tree_pred))
print(confusion_matrix(y_test, tree_pred))
print(classification_report(y_test, tree_pred))
```

```
Decision Tree Accuracy: 0.7755102040816326
[[219  36]
 [ 30   9]]
      precision    recall   f1-score   support
          0       0.88      0.86      0.87      255
          1       0.20      0.23      0.21       39

   accuracy                           0.78      294
macro avg       0.54      0.54      0.54      294
weighted avg    0.79      0.78      0.78      294
```

```
In [37]: log_acc = accuracy_score(y_test, log_pred)
tree_acc = accuracy_score(y_test, tree_pred)
```


```
In [38]: log_acc
```

```
Out[38]: 0.8945578231292517
```

```
In [39]: tree_acc
```

```
Out[39]: 0.7755102040816326
```

```
In [40]: sns.countplot(x=df['Attrition'])
plt.show()
```



```
In [41]: sns.countplot(x=df['Department'], hue=df['Attrition'])
plt.xticks(rotation=90)
plt.show()
```

